

▼ #ELEN 6885 Reinforcement Learning Coding Assignment (Part 4)# There are a lot of official and unofficial tutorials about Tensorflow, and there are also many open-source projects written in Tensorflow. You can refer to those resources according to your interest. In this part of homework 4, only knowledge of Deep Reinforcement Learning and basic programming skills will be needed.

Please put your code into the block marked by\ #####\ # YOUR CODE STARTS HERE\ # YOUR CODE ENDS HERE\ #####\ Normally you don't need to edit anything outside of the block. If you do want to edit something, please use a similar manner to mark you edits.

```

In [0]: 1 import numpy as np
2 import tensorflow as tf
3
4 # DQN
5 class DQN:
6     def __init__(
7         self,
8         actions_num,
9         state_size,
10        learning_rate = 0.001,
11        gamma = 0.99,
12        epsilon_min = 0.05,
13        epsilon_start = 0.9,
14        replace_target_iter = 300,
15        memory_size = 500,
16        batch_size = 2,
17        epsilon_increment = None,
18    ):
19        self.actions_num = actions_num
20        self.state_size = state_size
21        self.lr = learning_rate
22        self.gamma = gamma
23        self.epsilon_min = epsilon_min
24        self.replace_target_iter = replace_target_iter
25        self.memory_size = memory_size
26        self.batch_size = batch_size
27        self.epsilon_increment = epsilon_increment
28        self.epsilon = epsilon_start if epsilon_increment is not None else epsilon_min
29        self.save_model_path = './weights/DQN_model.ckpt'
30        self.memory_counter = 0
31
32        # learned steps counter
33        self.steps_counter = 0
34
35        # initialize memory [s, a, r, s_, done]
36        self.memory = np.zeros((self.memory_size, state_size * 2 + 3))
37
38        # build target_net and q_net
39        self.build_net()
40        t_params = tf.get_collection('target_net_params')
41        q_params = tf.get_collection('q_net_params')
42        self.replace_target = [tf.assign(t, q) for t, q in zip(t_params, q_params)]
43
44        # gpu setting
45        config = tf.ConfigProto(log_device_placement=False, allow_soft_placement=True)
46        config.gpu_options.per_process_gpu_memory_fraction = 0.6
47        self.sess = tf.Session(config=config)
48
49        self.sess.run(tf.global_variables_initializer())
50
51    def build_net(self):
52        # build q_net
53        self.state = tf.placeholder(tf.float32, [None, self.state_size], name='state')
54        self.q_target = tf.placeholder(tf.float32, [None, self.actions_num], name='q_target')
55        with tf.variable_scope('q_net'):
56            # c_names(collections_names) are the collections to store variables

```

```

57     c_names, neurons_layer_1, w_initializer, b_initializer = \
58         ['q_net_params', tf.GraphKeys.GLOBAL_VARIABLES], 100, \
59         tf.random_normal_initializer(0., 0.3), tf.constant_initializer(0.)
60
61     # layer 1
62     with tf.variable_scope('layer_1'):
63         w_layer_1 = tf.get_variable('w_layer_1', [self.state_size, neurons_layer_1],
64                                     w_initializer, b_initializer)
65         b_layer_1 = tf.get_variable('b_layer_1', [1, neurons_layer_1],
66                                     b_initializer, b_initializer)
67         layer_1 = tf.nn.relu(tf.matmul(self.state, w_layer_1) + b_layer_1)
68
69     # layer 2
70     with tf.variable_scope('layer_2'):
71         w_layer_2 = tf.get_variable('w_layer_2', [neurons_layer_1, self.actions_num],
72                                     w_initializer, b_initializer)
73         b_layer_2 = tf.get_variable('b_layer_2', [1, self.actions_num],
74                                     b_initializer, b_initializer)
75         self.q_value = tf.matmul(layer_1, w_layer_2) + b_layer_2
76
77     with tf.variable_scope('loss'):
78         self.loss = tf.reduce_mean(tf.squared_difference(self.q_target, self.q_value))
79
80     with tf.variable_scope('train'):
81         self._train_op = tf.train.AdamOptimizer(self.lr).minimize(self.loss)
82
83     # build target_net
84     self.state_t = tf.placeholder(tf.float32, [None, self.state_size])
85     with tf.variable_scope('target_net'):
86         # c_names(collections_names) are the collections to store variables
87         c_names = ['target_net_params', tf.GraphKeys.GLOBAL_VARIABLES]
88
89         # layer 1
90         with tf.variable_scope('layer_1'):
91             w_layer_1 = tf.get_variable('w_layer_1', [self.state_size, neurons_layer_1],
92                                         w_initializer, b_initializer)
93             b_layer_1 = tf.get_variable('b_layer_1', [1, neurons_layer_1],
94                                         b_initializer, b_initializer)
95             layer_1 = tf.nn.relu(tf.matmul(self.state_t, w_layer_1) + b_layer_1)
96
97         # layer 2
98         with tf.variable_scope('layer_2'):
99             w_layer_2 = tf.get_variable('w_layer_2', [neurons_layer_1, self.actions_num],
100                                         w_initializer, b_initializer)
101             b_layer_2 = tf.get_variable('b_layer_2', [1, self.actions_num],
102                                         b_initializer, b_initializer)
103             self.q_next = tf.matmul(layer_1, w_layer_2) + b_layer_2
104
105     # YOUR CODE ENDS HERE
106     #####
107
108     def store_transition(self, s, a, r, s_, done):
109         s=s.reshape(-1)
110         s_=s_.reshape(-1)
111         transition = np.hstack((s, [a, r], s_, done))
112         # replace the old memory with new observations
113         index = self.memory_counter % self.memory_size
114         self.memory[index, :] = transition
115
116         self.memory_counter += 1
117
118     def choose_action(self, observation):
119         # to have batch dimension when fed into tf placeholder

```

```

114 observation = observation[np.newaxis, :]
115 # epsilon-greedy
116 ▼ if np.random.uniform() > self.epsilon:
117     action_values = self.sess.run(self.q_value, feed_dict={self.state: batch_memory})
118     action = np.argmax(action_values)
119 ▼ else:
120     action = np.random.randint(0, self.actions_num)
121     return action
122
123 ▼ def learn(self):
124     # replace target parameters every once a while
125     ▼ if self.steps_counter % self.replace_target_iter == 0:
126         self.sess.run(self.replace_target)
127
128     # sample a batch from the memory
129     ▼ if self.memory_counter > self.memory_size:
130         sample_index = np.random.choice(self.memory_size, size=self.batch_size)
131     ▼ else:
132         sample_index = np.random.choice(self.memory_counter, size=self.batch_size)
133     batch_memory = self.memory[sample_index, :]
134
135     q_next, q_value = self.sess.run(
136         [self.q_next, self.q_value],
137     ▼ feed_dict={
138         self.state_t: batch_memory[:, -self.state_size+1:-1], # fixed parameters
139         self.state: batch_memory[:, :self.state_size], # newest parameters
140     })
141
142     # calculate q_target
143     q_target = q_value.copy()
144
145
146     # only change the action-values of this batch, because we only can observe one action
147     batch_index = np.arange(self.batch_size, dtype=np.int32)
148     act_index = batch_memory[:, self.state_size].astype(int)
149     reward = batch_memory[:, self.state_size + 1]
150     done = batch_memory[:, -1]
151     #####
152     # YOUR CODE STARTS HERE
153     q_target[batch_index, act_index] = reward + self.gamma * np.max(q_value, axis=-1)
154     ▼ for i in range(done.shape[0]):
155     ▼     if done[i] == 1.0:
156         q_target[i, :] = reward[i]
157
158     # YOUR CODE ENDS HERE
159     #####
160
161     # train q_net
162     ▼ _, self.cost = self.sess.run([self._train_op, self.loss],
163     ▼ feed_dict={self.state: batch_memory,
164                 self.q_target: q_target})
165
166     # change epsilon
167     self.epsilon = self.epsilon - self.epsilon_increment if self.epsilon > self.epsilon_min else self.epsilon
168     self.steps_counter += 1
169
170 ▼ def store(self):
171     saver = tf.train.Saver()

```

```
171     saver.save(self.sess, self.save_model_path)
172
173     def restore(self):
174         saver = tf.train.Saver()
175         saver.restore(self.sess, self.save_model_path)
176
177
178
179
180
```

Type *Markdown* and LaTeX: α^2

```
In [0]: 1  import gym
        2  # cart pole gym environment
        3  env = gym.make("CartPole-v0")
        4  env._max_episode_steps = 500
        5  # state and action space
        6  print(env.action_space)
        7  print(env.observation_space)
        8  # observation
        9  env.reset()
       10  # state, reward, done, info
       11  print(env.step(1))
```

Discrete(2)

Box(4,)

(array([0.03984107, 0.17240988, -0.04123799, -0.2650715]), 1.0, False, {})

```

In [0]: 1 ▾ # play the game and train the network
2 np.set_printoptions(threshold=np.inf)
3 episode_length_set = []
4 tf.reset_default_graph()
5 total_time_steps = 100000
6
7 ▾ RL = DQN(actions_num = 2, gamma = 0.99,
8           state_size = 4, epsilon_start = 1,
9           learning_rate = 1e-3, epsilon_min = 0.01,
10          replace_target_iter = 100, memory_size = 5000,
11          epsilon_increment = 0.00001,)
12
13 new_state = env.reset()
14 done = False
15 episode_length_counter = 0
16 ▾ for step in range(total_time_steps):
17     #####
18     # YOUR CODE STARTS HERE
19     action = RL.choose_action(new_state)
20     next_state, reward, done, _ = env.step(action)
21     RL.store_transition(new_state, action, reward, next_state, done)
22     ▾ if done:
23         new_state = env.reset()
24         episode_length_set.append(episode_length_counter)
25         episode_length_counter = 0
26     ▾ else:
27         new_state = next_state
28
29     # YOUR CODE ENDS HERE
30     #####
31     #print(step)
32     ▾ if step > 200:
33         RL.learn()
34         episode_length_counter += 1
35     ▾ if episode_length_counter == 500:
36         print('we hit 500')
37         RL.store()

```


```

we hit 500
we hit 500
we hit 500
we hit 500
we hit 500
we hit 500
we hit 500
we hit 500
we hit 500

```

```
In [0]: 1 from matplotlib import pyplot as plt  
2 plt.plot(episode_length_set)
```

Out[146]: [<matplotlib.lines.Line2D at 0x7f5fc28c7978>]



```

In [0]: 1 # test our network
2 tf.reset_default_graph()
3 RL = DQN(actions_num = 2, gamma = 1,
4           state_size = 4, epsilon_start = 1,
5           learning_rate = 1e-3, epsilon_min = 0,
6           replace_target_iter = 100, memory_size = 5000,
7           epsilon_increment = None,)
8 # load saved parameters
9 RL.restore()
10 # run 100 trails and print how long can the agent hold the cart pole f
11 length_counter = []
12 for i in range(100):
13     #####
14     # YOUR CODE STARTS HERE
15     new_state = env.reset()
16     episode_length = 0
17     done = False
18     while not done:
19         action = RL.choose_action(new_state)
20         next_state, reward, done, _ = env.step(action)
21         RL.store_transition(new_state, action, reward, next_state, done)
22         episode_length += 1
23     if done:
24         new_state = env.reset()
25         length_counter.append(episode_length)
26     else:
27         new_state = next_state
28
29     # YOUR CODE ENDS HERE
30     #####
31

```


INFO:tensorflow:Restoring parameters from ./weights/DQN_model.ckpt

```

In [0]: 1 plt.plot(length_counter)

```

Out[148]: [<matplotlib.lines.Line2D at 0x7f5fc1533a20>]




```
In [0]: 1 print(length_counter)
```

```
[495, 500, 481, 500, 500, 500, 500, 500, 500, 500, 469, 500, 500, 500, 500, 500, 500, 500, 496, 500, 500, 500, 500, 500, 500, 500, 500, 500, 500, 500, 500, 500, 454, 500, 500, 469, 500, 488, 500, 500, 500, 500, 500, 500, 500, 477, 500, 500, 500, 452, 500, 500, 427, 500, 500, 500, 500, 500, 500, 500, 500, 500, 500, 475, 500, 473, 410, 500, 500, 500, 500, 500, 500, 500, 500, 500, 445, 500, 500, 500, 500, 500, 500, 487, 500, 484, 500, 500, 500, 500, 500, 442, 500, 500, 500, 500, 500, 500, 500, 500, 500, 500, 481, 500, 500, 500]
```

You may find that the episode length doesn't stably improve as more training time is given. You can read chapter 3.2 of this paper <https://arxiv.org/pdf/1711.07478.pdf> (<https://arxiv.org/pdf/1711.07478.pdf>) if you are interested.